

Machine Learning 1

Muhammad Tariq

Table of contents


First up kmean()	1
Add more clusters to the plots	3
Hierarchical Clustering	4
Principal Component Analysis (PCA)	6
PCA TO THE RESUE	9
Interpreting PCA results	9

Today we explore the use of different data presentation tools on R like PCA

First up kmean()

Demo of using kmean() function in base R. First make some data with a known structure.

```
tmp <- c( rnorm(30,-3), rnorm(30,3))  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



Now we have some made up data in 'x' lets see how `kmeans()` works with this data

```
k <- kmeans(x,centers = 2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.235430	2.739656
2	2.739656	-3.235430

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:


```
[1] 41.05111 41.05111
(between_SS / total_SS = 92.9 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
```


```
[6] "betweenss"      "size"           "iter"           "ifault"
```

```
plot(x, col=k$cluster)
points(k$centers,col="blue",pch=15)
```



Add more clusters to the plots

```
k <- kmeans(x,centers = 4, nstart = 20)
plot(x, col=k$cluster)
points(k$centers,col="blue",pch=15)
```



key-point; K-means clustering is super popular but can be misused. one big limitation is that it can impose a clustering pattern on your data even if clear natural grouping doesn't exist- i.e it does what you tell it to do in terms of 'center'.

Hierarchical Clustering

the main function in “base” R for hierarchical clustering is called ‘hclust’.

You can just pass our dataset as is into ‘hclust()’ you must give “distance matrix” as input. We can get this from the ‘dist()’ function in R.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

The results of 'hclust)' dont have a useful 'print()' method but do have a speacial 'plot()' method.

```
plot(hc)
abline(h=8, col="red")
```



```
hclust (*, "complete")
```

To get our main cluster assignmnet (membership vector) we need to “cut” the tree at the big goal posts


```
grps <- cutree(hc, h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
table(grps)
```

```
grps
  1  2
30 30
```

```
plot(x)
```



Hierarchical clustering is distinct in that the dendrogram (tree figure) can reveal the potential grouping in your data (unlike K-means)

Principal Component Analysis (PCA)

PCA is a common and highly useful dimensionality reduction technique used in many fields - particularly bioinformatics.

Here we will analyze some data from UK on food consumption.

```
###data import
```

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)

head(x)
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267

3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139


```
rownames(x)<- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
x <- read.csv(url, row.names=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



one conventional plot that can be useful is called a “paris” plot.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```

PCA TO THE RESUE

The main function in base R for PCA is called 'prcomp()'

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Interpreting PCA results

The 'prcomp()' function returns a list object of our results with five attributes/components

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```


The two main “results” in here are ‘*pca\$x*’ and ‘*pca\$rotation*’. the first of these (*pca\$x*) contains the scores of the data on the new PC axis - we use these to make our “PCA plot”.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

```
library(ggplot2)


#Make a plot of pca$x with PC1 vs PC2
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point() +
  geom_text()
```



! PC1 enlightens the differences in the rows, in how different the foods are consumed.

The second major result is contained in the 'pca\$rotation' object or component. Let's plot this to see what PCA is picking up...

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```



lit shows how which food is eaten so differently.